

Seth Rojas

skrojas@arizona.edu | (520) 307-4663 | sethrojas.com | github.com/sethrojas21 | linkedin.com/in/skrojas/

EDUCATION

University of Arizona, W.A Franke Honors College

Bachelor of Science in Computer Science | Minors in Economics & Spanish

Tucson, Arizona

Expected May 2028

GPA: 4.0

Awards: Flinn Scholar, Hispanic Scholarship Fund Scholar, Hispanic Heritage Foundation: Silver in Technology, College Board National Hispanic Recognition Program Scholar, Arizona Seal of Biliteracy: Spanish, Dean's List with Distinction

Organizations: ColorStack, Society of Hispanic Engineers, CS Ambassadors, Intramural Basketball Champion

SKILLS

Programming languages: Java, Python, TypeScript, JavaScript, SQL, R, Swift, Assembly, C, HTML/CSS, XML

Technologies: Git, GitHub, React, Claude, GitHub Copilot, IBM Bob, scikit-learn, REST APIs, FastAPI, PostgreSQL, SQLite, SwiftData, Turso, pandas, NumPy, Matplotlib, Docker, Tidyverse, Google Cloud, Supabase

Courses: Data Structures & Algorithms, Object-Oriented Programming, Machine Learning, Systems Programming

WORK EXPERIENCE

IBM Storage CTO's Office

Tucson, Arizona

Software Development Engineer Intern

May 2026 – Present

- Developed a Python agent for demonstration at NVIDIA GTC Berlin and deployed to 50+ IBM sales reps, automating the process to find a customer's optimal storage infrastructure for AI inference from weeks to hours
- Quickly learned agentic AI in 2 weeks and architected a pipeline to process natural language into parallel simulation runs with results in under 7 minutes, providing data-driven recommendations using LLMs
- Developed modular unit tests for each component of the agent pipeline ensuring effectiveness before production
- Leveraged AI tools like IBM Bob to build weekly prototypes and incorporate feedback from 2 senior engineers
- Built working knowledge of modern AI inference infrastructure through research and presentations on KV-cache offloading, inference engines, NVIDIA Dynamo, and Red Hat OpenShift
- Collaborated with 3 IBM researchers to use their open-source software and co-author a published [blog](#)

University of Arizona Athletics

Tucson, Arizona

Data Science Intern

October 2025 – Present

- Deployed 4 machine learning models trained on 100+ features across 8 seasons to estimate a player's on-field contribution, supporting Arizona Football's recruiting and financial decision-making for 21 transfer portal recruits
- Built a full-stack Streamlit dashboard enabling scouts to filter, compare, and interpret player analytics through interactive charts and summaries, reducing the search for high-impact prospects at each position by up to 6x
- Architected a scalable data schema and ETL pipeline to ingest 200,000+ players into an Amazon Aurora PostgreSQL database, building Python loaders to boost team data analysis and optimize ML workflows
- Collaborated with general managers, football scouts, and financial directors, effectively communicating complex analytics through slideshow presentations, email summaries, and face-to-face conversations

University of Arizona Department of Computer Science

Tucson, Arizona

Independent Student Research Intern

January 2025 – October 2025

- Utilized strong analytical and problem-solving skills to create 2 novel basketball recruiting metrics that measure a player's team fit and on-court value, identifying successful recruits with 50.56% greater accuracy than chance
- Curated a 35,000+ player SQLite database and ETL pipeline for data analysis and machine learning, scraping 363 team sites with BeautifulSoup, extracting player data using LLM APIs, and integrating with 4 online databases
- Built and deployed a full-stack React dashboard with a FastAPI backend hosted on Google Cloud, communicating complex research findings through interactive visualizations for a 50+ non-technical audience

PROJECTS

Real-time Collaboration Spreadsheet (Class Project)

- Developed a Java client-server collaboration system using TCP sockets and object serialization to synchronize spreadsheet edits, formatting, images, chat, and active-cell presence for up to 5 concurrent users
- Collaborated with 3 peers to implement an MVC architecture, separating spreadsheet data, JavaFX interfaces, and collaboration logic to keep the application modular as its feature set expanded

UA AI Studio - HackArizona 2026

- Owned the end to end development of a RAG-powered AI teaching assistant that ingests classroom notes to deliver context-aware tutoring responses, using FastAPI for the backend and Supabase to store vector embeddings